

AutoML challenge winners

Victor Kochegarov

March 21, 2016

The most frequent prize-winners

djajetic (Damir Jajetic (Croatia))

4 out of 8 prizes: Final2 (2nd), AutoML3 (only one), Final3 (2nd), AutoML4 (2nd)

aad_freiburg Frank Hutter and collaborators (Freiburg, Germany)

7 out of 8 prizes: Final0 (3rd), AutoML1 (1st), Final1 (1st), AutoML2 (2nd), Final2 (3rd), Final3 (1st), AutoML4 (1st)

The most frequent prize-winners

djajetic (Damir Jajetic (Croatia))

4 out of 8 prizes: Final2 (2nd), AutoML3 (only one), Final3 (2nd), AutoML4 (2nd)

aad_freiburg Frank Hutter and collaborators (Freiburg, Germany)

7 out of 8 prizes: Final0 (3rd), AutoML1 (1st), Final1 (1st), AutoML2 (2nd), Final2 (3rd), Final3 (1st), AutoML4 (1st)

ideal.intel.analytics

4 out of 4 code-free prizes: Final0 (1st), Final1 (2nd), Final2 (2nd), Final3 (3rd)

The most frequent prize-winners

djajetic (Damir Jajetic (Croatia))

4 out of 8 prizes: Final2 (2nd), AutoML3 (only one), Final3 (2nd), AutoML4 (2nd)

aad_freiburg Frank Hutter and collaborators (Freiburg, Germany)

7 out of 8 prizes: Final0 (3rd), AutoML1 (1st), Final1 (1st), AutoML2 (2nd), Final2 (3rd), Final3 (1st), AutoML4 (1st)

ideal.intel.analytics

4 out of 4 code-free prizes: Final0 (1st), Final1 (2nd), Final2 (2nd), Final3 (3rd)

djajetic approach

Simple cross validation through models


linear__model.LogisticRegression
naive__bayes.GaussianNB
f16.RandomForestClassifier
ensemble.GradientBoostingClassifier
f16.ExtraTreesClassifier
ensemble.AdaBoostClassifier
ensemble.AdaBoostClassifier (base__estimator=f16.RF)
neighbors.KNeighborsClassifier

djajetic approach

```
models = [
    {"model": 'linear_model.LogisticRegression(random_state=1)',
     "blend_group": "LC",
     "getter": "lestimators = model_last.get_params()['penalty']",
     "updater": "tries_left = 0",
     "setter": "return in updater will end process",
     "generator": "penalty='l2', dual=False, C=1.0 @@ " \
                  "penalty='l1', dual=False, C=1.0 @@ " \
                  "penalty='l2', dual=False, C=2.0 @@ " \
                  "penalty='l2', dual=False, C=0.5 @@ " \
                  "penalty='l2', dual=True, C=1.0 @@ " \
                  "penalty='l2', dual=True, C=0.5 @@ " \
                  "penalty='l2', dual=False, C=4.0 @@ " \
                  "penalty='l2', dual=False, C=8.0 @@ " \
                  "penalty='l2', dual=False, C=1.0 @@ " \
                  "penalty='l1', dual=False, C=1.0 @@ " \
                  "penalty='l2', dual=False, C=2.0 @@ " \
                  "penalty='l2', dual=False, C=0.5 @@ " \
                  "penalty='l2', dual=True, C=1.0 @@ " \
                  "penalty='l2', dual=True, C=0.5 @@ " \
                  "penalty='l2', dual=False, C=4.0 @@ " \
                  "penalty='l2', dual=False, C=8.0 @@ "
    },
    {"model": 'naive_bayes.GaussianNB()',
     "blend_group": "NB",
     "getter": "lestimators = model_last.get_params()",
     "updater": "tries_left = 0",
     "setter": "",
     "generator": ""
    },
]
```

aad_freiburg approach

State-of-the-art before this guy (Auto-WEKA):



Data preprocessing

Algorithms that change data values

- 1 Rescaling of inputs
- 2 Imputation of missing values
- 3 Balancing of the target classes

Feature preprocessing

Algorithms that don't change data, but reduce number of features

- 1 PCA
- 2 Linear SVM (non zero model coefficients)
- 3 Feature agglomeration
- 4 Random forest
- 5 Extremely randomized forest
- 6 ...

Feature preprocessing

Well-established classification algorithms

- 1 General linear models
- 2 SVM
- 3 Discriminant analysis
- 4 Nearest neighbors
- 5 Naive Bayes
- 6 Decision trees
- 7 Ensemble methods


Auto-sklearn features result list

name	# λ	cat (cond)	cont (cond)
AdaBoost (AB)	3	-	3 (-)
Bernoulli naïve Bayes	2	1 (-)	1 (-)
decision tree (DT)	3	1 (-)	2 (-)
extreml. rand. trees	5	2 (-)	3 (-)
Gaussian naïve Bayes	-	-	-
gradient boosting (GB)	6	-	6 (-)
kNN	3	2 (-)	1 (-)
LDA	2	-	2 (-)
linear SVM	5	3 (-)	2 (-)
kernel SVM	8	3 (-)	5 (2)
multinomial naïve Bayes	2	1 (-)	1 (-)
passive aggressive	3	1 (-)	2 (-)
QDA	2	-	2 (-)
random forest (RF)	5	2 (-)	3 (-)
ridge regression (RR)	2	-	2 (-)
SGD	9	3 (-)	6 (3)

name	# λ	cat (cond)	cont (cond)
extreml. rand. trees prepr.	5	2 (-)	3 (-)
fast ICA	4	3 (-)	1 (-)
feature agglomeration	3	2 (-)	1 (-)
kernel PCA	5	1 (-)	4 (3)
rand. kitchen sinks	2	-	2 (-)
linear SVM prepr.	5	3 (-)	2 (-)
no preprocessing	-	-	-
nystroem sampler	5	1 (-)	4 (3)
PCA	2	1 (-)	1 (-)
random trees embed.	4	-	4 (-)
select percentile	2	1 (-)	1 (-)
select rates	3	2 (-)	1 (-)
imputation	1	1 (-)	-
balancing	1	1 (-)	-
rescaling	1	1 (-)	-

aad_freiburg approach

This guy added two more stages:



Meta-Learning

This is used to warm-start Bayesian optimizer

- 1 Given 140 datasets, compute it's metafeatures (total 38)
 - Simple features (N samples, N features, N classes etc)

Meta-Learning

This is used to warm-start Bayesian optimizer

- 1 Given 140 datasets, compute it's metafeatures (total 38)
 - Simple features (N samples, N features, N classes etc)
 - Information-theoretic features (entropy of features, Entropy of classes, feature noisiness etc)

Meta-Learning

This is used to warm-start Bayesian optimizer

- 1 Given 140 datasets, compute it's metafeatures (total 38)
 - Simple features (N samples, N features, N classes etc)
 - Information-theoretic features (entropy of features, Entropy of classes, feature noisiness etc)
 - Statistical metafeatures (Departure from normality, univariate skewness and kurtosis etc)

Meta-Learning

This is used to warm-start Bayesian optimizer

- ① Given 140 datasets, compute it's metafeatures (total 38)
 - Simple features (N samples, N features, N classes etc)
 - Information-theoretic features (entropy of features, Entropy of classes, feature noisiness etc)
 - Statistical metafeatures (Departure from normality, univariate skewness and kurtosis etc)
- ② Using Bayesian optimizer for each dataset choose ML framework with strongest performance

Meta-Learning

This is used to warm-start Bayesian optimizer

- ① Given 140 datasets, compute it's metafeatures (total 38)
 - Simple features (N samples, N features, N classes etc)
 - Information-theoretic features (entropy of features, Entropy of classes, feature noisiness etc)
 - Statistical metafeatures (Departure from normality, univariate skewness and kurtosis etc)
- ② Using Bayesian optimizer for each dataset choose ML framework with strongest performance
- ③ For a new dataset D find k nearest from the given 140 ones (L_1 -metric on metafeatures dataset).

Meta-Learning

This is used to warm-start Bayesian optimizer



- ① Given 140 datasets, compute it's metafeatures (total 38)
 - Simple features (N samples, N features, N classes etc)
 - Information-theoretic features (entropy of features, Entropy of classes, feature noisiness etc)
 - Statistical metafeatures (Departure from normality, univariate skewness and kurtosis etc)
- ② Using Bayesian optimizer for each dataset choose ML framework with strongest performance
- ③ For a new dataset D find k nearest from the given 140 ones (L_1 -metric on metafeatures dataset).
- ④ Warm-start Bayesian optimizer with k found parameters

Meta-Learning

This is used to warm-start Bayesian optimizer

- ① Given 140 datasets, compute it's metafeatures (total 38)
 - Simple features (N samples, N features, N classes etc)
 - Information-theoretic features (entropy of features, Entropy of classes, feature noisiness etc)
 - Statistical metafeatures (Departure from normality, univariate skewness and kurtosis etc)
- ② Using Bayesian optimizer for each dataset choose ML framework with strongest performance
- ③ For a new dataset D find k nearest from the given 140 ones (L_1 -metric on metafeatures dataset).
- ④ Warm-start Bayesian optimizer with k found parameters

Bayesian optimization reminder



Ensemble building

After Bayesian optimization one gets set of ML frameworks:

$$M = \{M_i\}$$

Algorithm:

Ensemble building

After Bayesian optimization one gets set of ML frameworks:

$$M = \{M_i\}$$

Algorithm:

- 1 Start with the empty ensemble

Ensemble building

After Bayesian optimization one gets set of ML frameworks:

$$M = \{M_i\}$$

Algorithm:

- 1 Start with the empty ensemble
- 2 Add to the ensemble the model in the library that maximizes the ensemble's performance to the error metric on a validation set

Ensemble building

After Bayesian optimization one gets set of ML frameworks:

$$M = \{M_i\}$$

Algorithm:

- 1 Start with the empty ensemble
- 2 Add to the ensemble the model in the library that maximizes the ensemble's performance to the error metric on a validation set
- 3 Repeat Step 2 for a fixed number of iterations or until all the models have been used

Ensemble building

After Bayesian optimization one gets set of ML frameworks:

$$M = \{M_i\}$$

Algorithm:

- 1 Start with the empty ensemble
- 2 Add to the ensemble the model in the library that maximizes the ensemble's performance to the error metric on a validation set
- 3 Repeat Step 2 for a fixed number of iterations or until all the models have been used
- 4 Return the ensemble from the nested set of ensembles that has maximum performance on the validation set

Ensemble building

After Bayesian optimization one gets set of ML frameworks:

$$M = \{M_i\}$$

Algorithm:

- 1 Start with the empty ensemble
- 2 Add to the ensemble the model in the library that maximizes the ensemble's performance to the error metric on a validation set
- 3 Repeat Step 2 for a fixed number of iterations or until all the models have been used
- 4 Return the ensemble from the nested set of ensembles that has maximum performance on the validation set

Ensemble building

After Bayesian optimization one gets set of ML frameworks:

$$M = \{M_i\}$$

Algorithm:

- 1 Start with the empty ensemble
- 2 Add to the ensemble the model in the library that maximizes the ensemble's performance to the error metric on a validation set
- 3 Repeat Step 2 for a fixed number of iterations or until all the models have been used
- 4 Return the ensemble from the nested set of ensembles that has maximum performance on the validation set

Ensemble building

With these improvements

- 1 Selection with replacement: adds model only if it increases performance

Ensemble building

With these improvements

- 1 Selection with replacement: adds model only if it increases performance
- 2 Sorted Ensemble Initialization: start with non empty ensemble

Ensemble building

With these improvements

- 1 Selection with replacement: adds model only if it increases performance
- 2 Sorted Ensemble Initialization: start with non empty ensemble
- 3 Bagged Ensemble Selection: repeat all the above several (20) times for bagged models

Ensemble building

With these improvements

- 1 Selection with replacement: adds model only if it increases performance
- 2 Sorted Ensemble Initialization: start with non empty ensemble
- 3 Bagged Ensemble Selection: repeat all the above several (20) times for bagged models

Ensemble building

With these improvements

- 1 Selection with replacement: adds model only if it increases performance
- 2 Sorted Ensemble Initialization: start with non empty ensemble
- 3 Bagged Ensemble Selection: repeat all the above several (20) times for bagged models

Auto-sklearn results

